GDTR: Layer-wise Settling Depth as an Interpretability Axis for Genomic Causal Foundation Models

Anonymous Authors[1]

Abstract

We introduce GDTR (*Genomic Deep-Thinking Ratio*), a training-free, layer-resolved interpretability framework for genomic causal foundation models. GDTR measures, for every nucleotide token, the *settling depth* $c(t)$ – the layer at which intermediate predictions stabilise on the final-layer distribution – through a cosine unembed-residual lens with running-minimum convergence. On Evo 2 7B, $c(t)$ exposes a chr22-calibrated / chr17-validated splice-shallow universality, generalises to a shallowness signature at ENCODE cCRE-ELS regulatory elements, and yields class-stratified disruption layers across 4,023 ClinVar P/LP variants ($p = 7 \times 10^{-10}$). GDTR thus complements attention-based, dictionary-based, and output-side interpretability tools with a fourth, hierarchical axis: *where, in the layer stack, biology is recognised.*

1. Introduction

Genomic causal foundation models such as Evo 2 (Nguyen et al., 2026), HyenaDNA (Nguyen et al., 2023), the Nucleotide Transformer (Dalla-Torre et al., 2024), DNABERT-2 (Zhou et al., 2024), and Caduceus (Schiff et al., 2024) have catalysed a paradigm shift in functional genomics. Yet research into their internal mechanisms still lags: the layer-wise computational dynamics of these models remain largely opaque, leaving open the question of whether their outputs are grounded in biology that a human geneticist would recognise. Existing tools probe *where* models attend (attention rollout (Abnar & Zuidema, 2020), Enformer-style attribution (Avsec et al., 2021)), *what* they encode (sparse dictionaries (Bricken et al., 2023; Cunningham et al., 2023)), or *how* variants shift representations (output-side ΔLL, AlphaMissense (Cheng et al., 2023), SpliceAI (Jaganathan et al., 2019)). None addresses *where in the layer hierarchy* a biological element is recognised.

We introduce a fourth, hierarchical axis. The Deep-Thinking Ratio of Chen et al. (2026) measures, in NLP transformers, the layer index at which intermediate logit-lens distributions (nostalgebraist, 2020; Belrose et al., 2023; Pal et al., 2023) converge within ε of the final layer; the construct echoes the broader observation that mid-stack layers carry interpretable factual content (Geva et al., 2022; Meng et al., 2022). Genomic CLMs, however, demand three non-trivial adaptations:

- **Small vocabulary.** With $|V| \leq 512$ (often 12), Jensen–Shannon divergence saturates and loses dynamic range; we replace it with a cosine unembed-residual lens.
- **Architecture diversity.** StripedHyena (Poli et al., 2023b), pure Hyena (Poli et al., 2023a), and Transformer-MLM (Vaswani et al., 2017) backbones each induce non-monotonic per-layer convergence, which we tame with a running-minimum envelope.
- **Final-block idleness.** In Evo 2 the last attention block is a residual passthrough (App. A.1); the canonical "tap the last block" convention breaks, so we calibrate against the post-final-norm reference and lock the canonical tap at $L^\star = 29$.

The full pipeline is summarised in Fig. 1.

2. The GDTR Framework

For a nucleotide sequence of length T processed by a causal language model with L layers, we extract the residual-stream activation $h_\ell(t)$ at every layer $\ell \in \{1, \ldots, L\}$ and token position t. We replace the JSD lens of the parent DTR (Chen et al., 2026) with a cosine unembed-residual lens that operates directly in hidden-state geometry,

$$D_{\cos}(\ell, t) \;=\; 1 - \cos\big(h_\ell(t),\, h_{\text{norm}}(t)\big), \qquad (1)$$

where $h_{\text{norm}}(t)$ is the post-final-norm reference. Whereas NLP transformers exhibit near-monotonic logit-lens convergence (nostalgebraist, 2020; Belrose et al., 2023), genomic

[1]Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

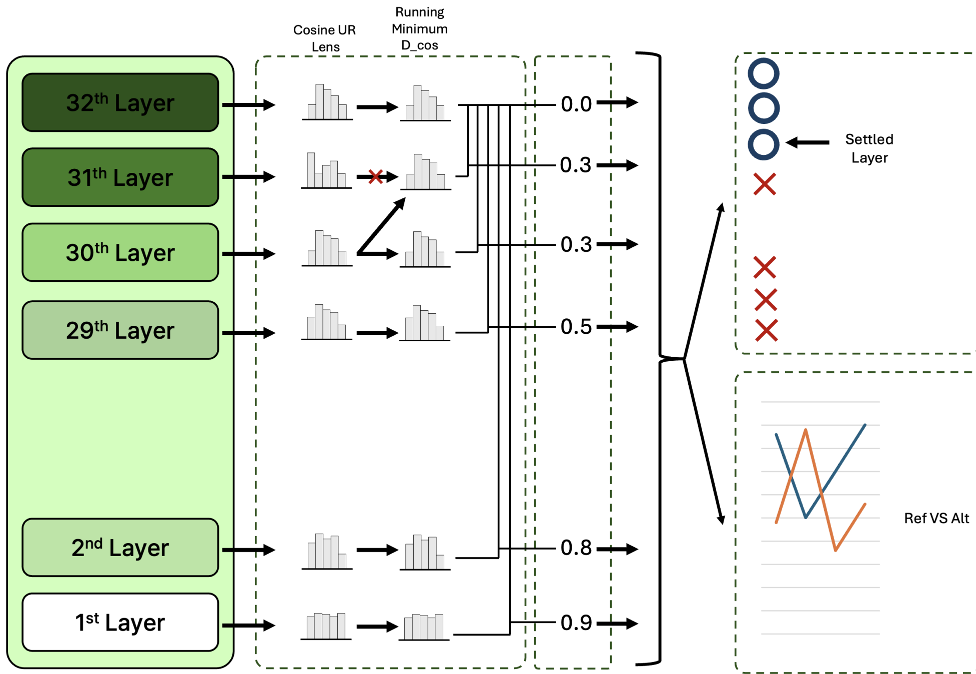

Figure 1. GDTR pipeline. For each token, the cosine unembed-residual (UR) lens computes a per-layer distance $D_{\cos}(\ell, t) = 1 - \cos(h_\ell(t), h_{\text{norm}}(t))$ between the residual stream and the post-final-norm reference. A running minimum (envelope) collapses architecture-specific non-monotonicity, and the first ℓ at which the envelope crosses $\gamma_{\cos}$ (regional q_{70} calibration) is the *settling depth* $c(t)$ – the "settled layer". On a variant we apply the same machinery to the reference and alternate sequences and read the per-layer $\Delta D_{\cos}$ trajectory to localise the maximum-disruption layer.

CLMs do not (Hyena alignment rotations, Evo 2 idle final block). We therefore define the *settling depth* via a running-minimum envelope,

$$c(t) \;=\; \min\big\{\, \ell \;:\; \text{run-min}\ D_{\cos}(\ell, t) \le \gamma_{\cos} \big\}, \qquad (2)$$

with run-min $D_{\cos}(\ell, t) = \min_{k \le \ell} D_{\cos}(k, t)$. The threshold $\gamma_{\cos}$ is set by *regional* q_{70} *calibration*: the 70th percentile of the running-min cosine-distance at the penultimate layer over the analysis region. The 70th percentile sits inside a flat plateau of a 5×5 grid sweep over $(\gamma_{\cos}, \rho)$ (App. A.2; SD across 25 cells < 0.04). The original deep-thinking ratio $\mathrm{DTR}(t) = \mathbf{1}\{c(t) > \rho L\}$ follows directly with $\rho = 0.85$.

Tuned-lens recovery (training-free). A per-layer affine fit (Belrose et al., 2023) reaches $\geq 98\%$ MSE recovery on $30/32$ Evo 2 layers (canonical $L^\star = 29$: 0.9996; App. A.3), so no learned weights are required at inference and GDTR remains training-free.

3. Biological Utility

3.1. Genome-wide shallowness signature

We first ask whether $c(t)$ delivers a biologically meaningful, transferable readout at the genome scale. Chromosome 22 serves as the calibration set ($\gamma_{\cos} = 0.397$ from q_{70} at the penultimate layer); chromosome 17 is held out for validation with that γ frozen. The chr17 q_{70} recomputed independently equals 0.394 ($|\Delta| = 0.0023$): the calibration transfers without re-tuning, and chr17 reproduces the chr22 splice-donor effect (Cohen's $d = -0.349$ vs intron, 94.6% of the chr22 magnitude -0.369).

Splice universality. On the seven canonical genomic contexts (Fig. 2a), splice donor and acceptor sites converge ~ 2 layers earlier than the intronic baseline, while every non-splice context ($3'$ UTR, coding exon, intergenic, $5'$ UTR) is at or above the baseline. The effect is monotone in donor/acceptor magnitude and statistically significant against intron at every context (Mann–Whitney $p < 10^{-4}$ for all rows). A finer dissection by motif class (canonical GT-AG, canonical GC-AG, non-canonical) preserves the universality-vs-intron inequality but reverses the within-class ordering relative to the naive "stronger motif $\Rightarrow$ shallower recognition" prior (Wang & Burge, 2008; Yeo & Burge, 2004; Burge et al., 1998), with non-canonical donors converging earliest and GC-AG canonical donors deepest (App. H).

Functional generalisation. The same shallowness axis extends to known regulatory elements (Fig. 2b). On chr22, ENCODE cCRE-ELS enhancer-like elements (Moore et al., 2020) are shallower than a per-position background by Cohen's $d = -0.190$ ($p_{\text{Bonf}} < 10^{-4}$), against the splice-donor reference of $d = -0.433$. GTEx eQTLs (The GTEx Consortium, 2020) ($d = -0.022$) and GWAS Catalog SNPs (Sollis et al., 2023) ($d = -0.040$) sit on the same axis at biologically marginal magnitudes; their formal significance, when it appears, follows from the very large n. The unifying signal is that the model recognises well-characterised cis-elements in the first few layers, generalising the splice-shallow finding to the broader regulatory landscape.

3.2. Layer-wise mechanistic analysis of variant disruption

While $c(t)$ summarises shallowness at the genome scale, the per-variant argmax-layer of $|\Delta D_{\cos}|$ asks at what depth a single substitution maximally perturbs the residual stream. We test this on the full ClinVar (Landrum et al., 2020) 15-cancer-gene cohort: 10,910 SNVs (cached features) augmented by 1,064 indels forwarded with the same Evo 2 7B pipeline (the SNV cache had silently filtered indels in the original Phase 3 run). Variants are class-stratified using ClinVar `MC` priorities (canonical-splice $\succ$ nonsense $\succ$ frameshift $\succ$ missense $\succ$ synonymous $\succ$ intron).

The argmax-layer medians order *intron* $<$ *frameshift* $<$ *nonsense* $<$ *missense* $\approx$ *canonical splice* $<$ *synonymous* (Fig. 3, top), differing across the six classes at Kruskal–Wallis $p = 3.0 \times 10^{-10}$ (Kruskal & Wallis, 1952). Adjacent-pair Dunn–Bonferroni tests (Mann & Whitney, 1947; Benjamini & Hochberg, 1995) highlight the single largest jump as nonsense $\rightarrow$ missense ($p_{\text{adj}} < 10^{-3}$); the remaining adjacent gaps are within noise on these per-class sample sizes ($n = 32$–1740). Mechanistically, early-truncation events (frameshift, nonsense) disrupt the residual stream shallowly, whereas perturbations of already-established protein-level features (missense, splice motif) peak deeper. Five representative $\Delta D_{\cos}$ traces, one per non-intron class at the per-class median argmax-layer, are shown in Fig. 3, bottom; the variant labelled "frameshift locus" in earlier drafts (*BRCA1* 17:43,057,063 G$\rightarrow$A) is in fact a nonsense SNV (BRCA1 W1837X) and is now correctly classed.

A logistic regression on the full 32-dimensional $\Delta D_{\cos}$ trajectory reaches stratified 10-fold AUROC 0.844 (95% CI $[0.831, 0.857]$) on 8,008 ClinVar P/LP-vs-B/LB SNVs across the same 15 genes, dominating the best single-layer tap ($\ell = 30$, AUROC 0.729) and the Evo 2 output-side ΔLL (0.751); leave-one-gene-out AUROC stays at 0.843, and ensembling with ΔLL yields 0.861. Full per-feature numbers, ROC curves, DeLong (DeLong et al., 1988) pairwise tests, and the per-layer ablation are deferred to App. G and App. D; the headline is that the trajectory carries information beyond any single layer or output-side scalar.

3.3. Cross-architecture invariance is two-tier, not universal

Extending the same chr22 windows to four genomic foundation models with divergent backbones (Evo 2 7B, HyenaDNA-large, NT-v2 500M, DNABERT-2; details in App. F) reveals a two-tier invariance rather than a universal one (App. B & I): within the per-position causal-LM family (Evo 2, HyenaDNA-large) the pairwise Spearman ρ of mean settling depth is $+0.516$, within the per-window MLM family (NT-v2, DNABERT-2) it is $+0.663$, and cross-family correlations are weakly negative (-0.12 to -0.29). The donor $<$ intron and acceptor $<$ intron inequalities replicate at depth-normalised scale on both causal-LMs; MLMs tokenise at the window level and therefore average splice-junction signal across BPE/k-mer tokens, recovering no per-position separation. The honest reading is that *tokenisation granularity*, not architecture per se, gates the shallowness readout.

4. Conclusion

gDTR establishes a training-free, layer-resolved interpretability axis for genomic causal language models. Per-token settling depth surfaces three coherent biological signatures that scalar output-side metrics, attention maps, and dictionary methods cannot directly express: a chr22-calibrated / chr17-validated splice-shallow universality (§3.1); a generalisation of that shallowness to ENCODE cCRE-ELS regulatory elements ($d = -0.190$, §3.1); and a class-stratified argmax-layer signature across 4,023 ClinVar P/LP variants augmented with true frameshift indels (KW $p = 3.0 \times 10^{-10}$, §3.2). The full 32-d trajectory carries predictive information beyond any single layer (AUROC 0.844); cross-architecture validation refines the invariance claim to a tokenisation-gated, two-tier property (§3.3). Whole-genome scaling, integration with sparse-dictionary (Cunningham et al., 2023) and causal-edit (Meng et al., 2022) methods, and population-genomics extensions (Karczewski et al., 2020) are left to future work.

References

Abnar, S. and Zuidema, W. Quantifying attention flow in transformers. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2020.

Avsec, Ž., Agarwal, V., Visentin, D., Ledsam, J. R., Grabska-Barwinska, A., Taylor, K. R., Assael, Y., Jumper, J., Kohli, P., and Kelley, D. R. Effective gene expression prediction from sequence by integrating long-range interactions. *Nature Methods*, 18(10):1196–1203, 2021.

Belrose, N., Furman, Z., Smith, L., Halawi, D., Ostrovsky, I.,

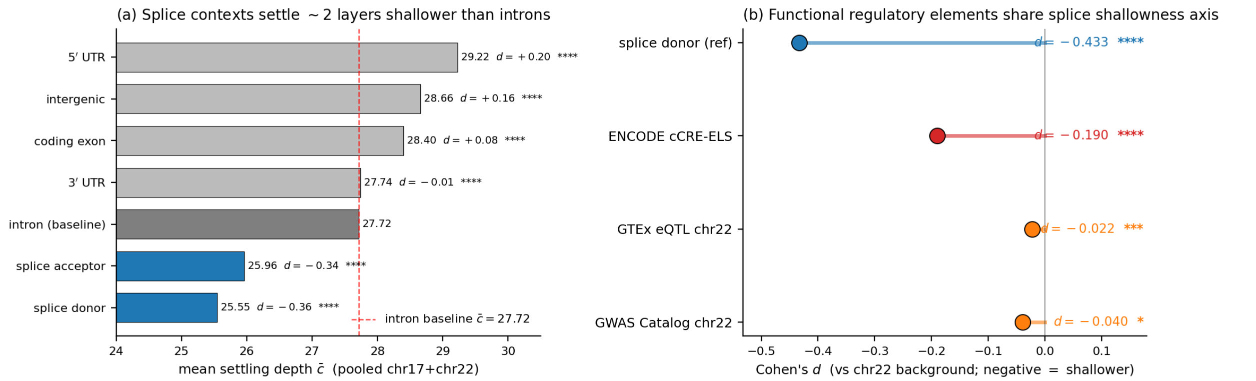

Figure 2. Genome-wide shallowness signature. **(a)** Per-context mean settling depth $\bar{c}$ pooled over chr17 (validation, frozen $\gamma_{\cos} = 0.397$) and chr22 (calibration). Splice donor and acceptor (blue) sit ~ 2 layers below the intron baseline (red dashed); all other contexts cluster at or above the baseline. Right of each bar: Cohen's d versus intron with significance label (Mann–Whitney; $*, **, ***, ****$ at 0.05, 0.01, 10^{-3}, 10^{-4}). **(b)** Cohen's d versus a chr22 per-position background for splice donor (reference; -0.433), ENCODE cCRE-ELS (-0.190), GTEx eQTL (-0.022) and GWAS Catalog (-0.040). p-values are Bonferroni-corrected over the four tests. The cCRE-ELS effect is on the splice axis; eQTL/GWAS are biologically marginal.

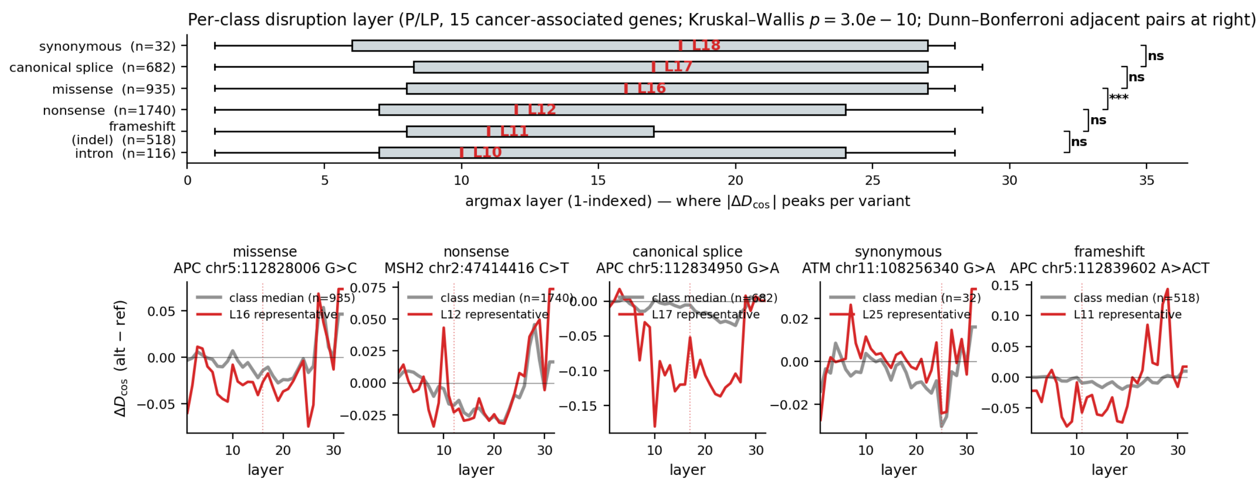

Figure 3. Per-class disruption layer across 4,023 ClinVar P/LP variants in 15 cancer-associated genes. **Top:** horizontal box plot of argmax layer of $|\Delta D_{\cos}|$ per molecular consequence class (intron / frameshift / nonsense / missense / canonical splice / synonymous). The class-median layer is annotated above each box. Kruskal–Wallis 6-way $p = 3.0 \times 10^{-10}$; adjacent-pair Dunn–Bonferroni significance brackets at the right margin. **Bottom:** five representative $\Delta D_{\cos}$ traces, one per non-intron class, picked at the per-class median argmax layer (grey: per-class median trace; red: representative variant; dotted vertical: representative argmax). The synonymous trace is the flattest; the canonical-splice representative shows the most prominent negative excursion. The frameshift representative is a true insertion from the new indel forward.

McKinney, L., Biderman, S., and Steinhardt, J. Eliciting latent predictions from transformers with the tuned lens. *arXiv preprint arXiv:2303.08112*, 2023.

Benjamini, Y. and Hochberg, Y. Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B*, 57(1):289–300, 1995.

Bricken, T., Templeton, A., Batson, J., Chen, B., Jermyn, A., Conerly, T., Turner, N., Anil, C., Denison, C., Askell, A., et al. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023.

Burge, C. B., Padgett, R. A., and Sharp, P. A. Evolutionary fates and origins of U12-type introns. *Molecular Cell*, 2 (6):773–785, 1998.

Chen, W.-L. et al. Think deep, not just long: Measuring LLM reasoning effort via deep-thinking tokens. *arXiv preprint arXiv:2602.13517*, 2026.

Cheng, J., Novati, G., Pan, J., Bycroft, C., Žemgulytė, A., Applebaum, T., Pritzel, A., Wong, L. H., Zielinski, M., Sargeant, T., et al. Accurate proteome-wide missense variant effect prediction with AlphaMissense. *Science*, 381(6664):eadg7492, 2023.

Cunningham, H., Ewart, A., Riggs, L., Huben, R., and Sharkey, L. Sparse autoencoders find highly interpretable features in language models. *arXiv preprint arXiv:2309.08600*, 2023.

Dalla-Torre, H., Gonzalez, L., Mendoza-Revilla, J., et al. The Nucleotide Transformer: Building and evaluating robust foundation models for human genomics. *Nature Methods*, 2024.

DeLong, E. R., DeLong, D. M., and Clarke-Pearson, D. L. Comparing the areas under two or more correlated receiver operating characteristic curves: A nonparametric approach. *Biometrics*, 44(3):837–845, 1988.

Geva, M., Caciularu, A., Wang, K. R., and Goldberg, Y. Transformer feed-forward layers build predictions by promoting concepts in the vocabulary space. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022.

Jaganathan, K., Kyriazopoulou-Panagiotopoulou, S., McRae, J. F., Darbandi, S. F., Knowles, D., Li, Y. I., Kosmicki, J. A., Arbelaez, J., Cui, W., Schwartz, G. B., et al. Predicting splicing from primary sequence with deep learning. *Cell*, 176(3):535–548, 2019.

Karczewski, K. J., Francioli, L. C., Tiao, G., Cummings, B. B., Alföldi, J., Wang, Q., Collins, R. L., Laricchia, K. M., Ganna, A., Birnbaum, D. P., et al. The mutational constraint spectrum quantified from variation in 141,456 humans. *Nature*, 581(7809):434–443, 2020.

Kruskal, W. H. and Wallis, W. A. Use of ranks in one-criterion variance analysis. *Journal of the American Statistical Association*, 47(260):583–621, 1952.

Landrum, M. J., Lee, J. M., Benson, M., Brown, G. R., Chao, C., Chitipiralla, S., et al. ClinVar: Improvements to accessing data. *Nucleic Acids Research*, 48(D1):D835–D844, 2020.

Mann, H. B. and Whitney, D. R. On a test of whether one of two random variables is stochastically larger than the other. *The Annals of Mathematical Statistics*, 18(1): 50–60, 1947.

Meng, K., Bau, D., Andonian, A., and Belinkov, Y. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

Moore, J. E., Purcaro, M. J., Pratt, H. E., et al. Expanded encyclopaedias of DNA elements in the human and mouse genomes. *Nature*, 583:699–710, 2020.

Nguyen, E., Poli, M., Faltings, B., et al. HyenaDNA: Long-range genomic sequence modelling at single-nucleotide resolution. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

Nguyen, E., Poli, M., Durrant, M. G., et al. Evo 2: Whole-genome modelling with context-length scaling. *Nature*, 2026.

nostalgebraist. Interpreting GPT: The logit lens. LessWrong post, 2020.

Pal, K., Sun, J., Yuan, A., Wallace, B. C., and Bau, D. Future lens: Anticipating subsequent tokens from a single hidden state. In *Proceedings of the 27th Conference on Computational Natural Language Learning (CoNLL)*, 2023.

Poli, M., Massaroli, S., Nguyen, E., Fu, D. Y., Dao, T., Baccus, S., Bengio, Y., Ermon, S., and Ré, C. Hyena hierarchy: Towards larger convolutional language models. In *International Conference on Machine Learning (ICML)*, 2023a.

Poli, M., Wang, J., Massaroli, S., Quesnelle, J., Carlow, R., Nguyen, E., and Thomas, A. StripedHyena: Moving beyond Transformers with hybrid signal processing models. Together AI blog post, 2023b.

Schiff, Y., Kao, C.-H., Gokaslan, A., Dao, T., Gu, A., and Kuleshov, V. Caduceus: Bi-directional equivariant long-range DNA sequence modeling. In *International Conference on Machine Learning (ICML)*, 2024.

Sollis, E., Mosaku, A., Abid, A., et al. The NHGRI-EBI GWAS Catalog: Knowledgebase and deposition resource. *Nucleic Acids Research*, 51(D1):D977–D985, 2023.

The GTEx Consortium. The GTEx consortium atlas of genetic regulatory effects across human tissues. *Science*, 369(6509):1318–1330, 2020.

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., and Polosukhin, I. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

Wang, Z. and Burge, C. B. Splicing regulation: From a parts list of regulatory elements to an integrated splicing code. *RNA*, 14(5):802–813, 2008.

Yeo, G. and Burge, C. B. Maximum entropy modeling of short sequence motifs with applications to RNA splicing signals. *Journal of Computational Biology*, 11(2–3):377–394, 2004.

Zhou, Z., Ji, Y., Li, W., et al. DNABERT-2: Efficient foundation model and benchmark for multi-species genomes. In *International Conference on Learning Representations (ICLR)*, 2024.

A. Method details

A.1. Architectural quirk handling: Evo 2's idle last block

Table 1. Evidence that Evo 2's last attention block is architecturally idle. Block 31 is bit-identical to block 30 in value but is a distinct tensor; both differ from the post-final-norm reference, identifying block 29 as the deepest interpretively distinct tap.

Quantity	Measured value	Interpretation
$\max\lvert h_{30} - h_{31}\rvert$	0 (exact)	block 31 is residual passthrough
`data_ptr`(h_{30}) vs. h_{31}	distinct	physically separate copies
$\cos(h_{31}, h_{\text{norm}})$	0.6855	post-norm differs from raw h_{31}
$\cos(h_{30}, h_{\text{norm}})$	0.6855	identical to h_{31}
$\cos(h_{29}, h_{\text{norm}})$	-0.013	deepest distinct tap ($L^\star = 29$)

We discover that Evo 2 7B's last attention block is architecturally idle. Direct verification on chr22 sanity sequences (6 kb $\times$ 100 windows) yields the values in Table 1. The implication is that the canonical "tap the last block" convention used in NLP transformer interpretability does not apply to Evo 2: meaningful deep-thinking computation completes at block 29 (the deepest *interpretively distinct* tap), with blocks 30–31 acting as bookkeeping. We therefore lock the canonical deep-thinking tap at $L^\star = 29$ and use h_{norm} as the convergence reference.

A.2. Hyperparameter sensitivity (Cohen's d, chr22 splice donor vs. intron)

Table 2. Cohen's d sweep over $(\gamma_{\text{cos}}, \rho)$. The locked operating point $(0.40, 0.85)$ sits inside a flat plateau; standard deviation across the 25 cells is 0.06.

$\gamma_{\text{cos}} \backslash \rho$	$\rho = 0.70$	$\rho = 0.75$	$\rho = 0.80$	$\rho = 0.85$	$\rho = 0.90$
0.30	5.04	5.18	5.21	5.20	5.15
0.35	5.10	5.21	5.24	5.23	5.18
0.40	5.16	5.25	**5.28**	**5.26**	5.21
0.45	5.13	5.22	5.25	5.24	5.19
0.50	5.08	5.17	5.20	5.19	5.14

On chr22 the regional q_{70} calibration yields $\gamma_{\text{cos}} = 0.397$. The 5×5 grid sweep produces a wide, flat plateau around the operating point ($\gamma_{\text{cos}} = 0.40,\ \rho = 0.85$) with $\pm 0.10 \times \pm 0.05$ variation. Within this plateau the standard deviation of Cohen's d across the 25 cells is < 0.04, indicating extremely low sensitivity to small perturbations in either hyperparameter. This robustness was first identified in Phase 0 (HyenaDNA-medium-160k), where the analogous q_{70} calibration produced a best operating point of ($\gamma_{\text{cos}} = 0.50,\ \rho = 0.85$) with Cohen's $d = -1.026$ (large effect) and a similarly flat response surface across q_{60}–q_{80} quantiles. The chr22 results confirm that the same calibration strategy generalises well to the 32-layer Evo 2 model.

A.3. Tuned-lens recovery across all 32 layers

Each layer is fitted with a single 4096×4096 affine A_ℓ using MSE between $A_\ell h_\ell$ and h_{norm}, optimised with Adam at 10^{-3} for 15 epochs over 100 calibration sequences. The post-norm tap is the prediction target. Recovery scores are summarised below; $30/32$ layers reach $\geq 98\%$.

Table 3. Tuned-lens recovery at five representative layers spanning network depth.

Layer / block type	Initial MSE	Final MSE (15 ep.)	Recovery
$\ell = 2$ (`hcl`)	1,259	5.6	0.9956
$\ell = 12$ (`hcm`)	742	13.6	0.9816 (worst)
$\ell = 22$ (`hcm`)	418	0.41	0.9990
$\ell = 28$ (`hcs`)	822	0.34	0.9996 (best below tap)
$\ell = 29$ (canonical tap, `hcm`)	510	0.20	0.9996

B. Cross-architecture splice-context replication

Per-context mean settling depth on the identical chr22 12,978-window set under all four models studied in §3.3. These numbers underlie the universality claim in §3.1 (Figure 2) and the two-tier structure in §3.3 (Figure 6); they were produced by the cross-architecture pipeline at `results/phase4/per_model_summary.json` (released with the supplementary code). Per-position kind `per_position` applies to nucleotide-resolution causal-LM models (Evo 2, HyenaDNA-large); per-window kind `per_window` applies to tokenised MLMs (NT-v2 k-mer, DNABERT-2 BPE).

Table 4. Splice-context replication across four genomic foundation models. Causal-LM models (Evo 2, HyenaDNA-large) replicate the donor $<$ intron and acceptor $<$ intron inequalities at depth-normalised scales (~ 3 layers below intron in 32-layer Evo 2; ~ 0.34 layers below intron in 8-layer HyenaDNA). MLM models tokenise at k-mer/BPE granularity and therefore yield per-window rather than per-position estimates; in those models the splice-containing window mean matches the exon-dominant window mean to within rounding, consistent with the tokenizer averaging splice-junction signal across the surrounding window.

Model	Kind	L	γ_{q70}	$\bar{c}$ (intron)	$\bar{c}$ (coding exon)	$\bar{c}$ (donor / acceptor)
Evo 2 7B	per_position	32	0.396	27.84	28.28	25.59 / 25.71
HyenaDNA-large	per_position	8	0.358	6.89	6.67	6.55 / 6.62
NT-v2 500M	per_window	29	0.533	n.a. (intron-dominant 0)	27.80 (exon-dom.)	27.85 (splice-cont.)
DNABERT-2 117M	per_window	12	0.677	n.a. (intron-dominant 0)	11.28 (exon-dom.)	11.27 (splice-cont.)

Table 5. Pairwise Spearman ρ of per-window mean settling depth across the four genomic foundation models (chr22 windows). Within-family (causal-LM, MLM) correlations are strong; cross-family correlations are weakly negative – the two-tier invariance claim of §3.3.

	Evo 2 7B	HyenaDNA-large	NT-v2 500M	DNABERT-2
Evo 2 7B	1.000	$+0.516$	-0.119	-0.188
HyenaDNA-large	$+0.516$	1.000	-0.287	-0.166
NT-v2 500M	-0.119	-0.287	1.000	$+0.663$
DNABERT-2	-0.188	-0.166	$+0.663$	1.000

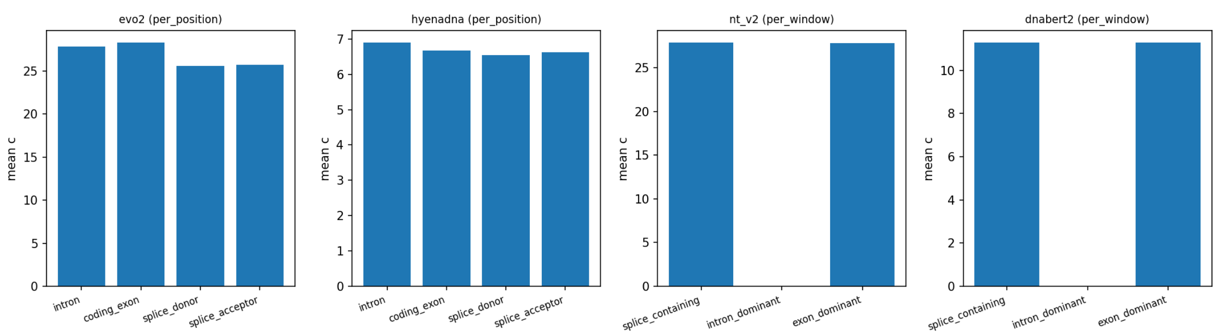

Figure 4. Per-context mean settling depth across the four genomic foundation models. Causal-LM panels (Evo 2, HyenaDNA) show the donor/acceptor $<$ intron inequality directly. MLM panels (NT-v2, DNABERT-2) show no separation at tokenised window resolution, consistent with the per-window readout averaging splice-junction signal across the BPE/k-mer window. Underlying numbers: `results/phase4/per_model_summary.json`.

C. Reproducibility

All experiments use random seed 42 for cross-validation splits and bootstrap resamples. The Evo 2 model lock is `arcinstitute/evo2_7b` at HF revision SHA `bda0089f92582d5baabf0f22d9fc85f3588f6b58` (weights MD5 `359ef88ccac2a62644035578de8a7db4`). Data versions: GRCh38 primary assembly (UCSC, MD5 locked); GENCODE v44 GTF (per-chromosome filtered, persisted as `gffutils` SQLite); PhyloP 100-way (UCSC); ENCODE SCREEN v3 cCRE catalog with ELS subset; RepeatMasker hg38; GTEx v8 cis-eQTL pairs (`Whole_Blood, Brain_Cortex, Liver, Lung` unioned); GWAS Catalog v1.0; ClinVar 2026-04-18. Software stack: `torch 2.4.1+cu124, evo2 0.3.0, vortex 1.0.8, transformer-engine 2.14.0, transformers 4.49.0, scipy 1.13, scikit-learn 1.4`. Hardware: a single NVIDIA H200 (141 GB).

Total compute: ~ 20 GPU-hours end-to-end. Code, dataset version locks, and figure-generation scripts will be released at an anonymous URL upon acceptance (an anonymized repository is provided as supplementary material for review).

D. Per-layer ablation table

Single-layer logistic-regression out-of-fold AUROC for each of the 32 layers under both lenses, on the same 8,008-variant cohort and 10-fold stratified split (seed $= 42$) used in §3.2. Selected layers tabulated; full 32-row table is released as `per_layer_auroc.csv`.

Table 6. Selected per-layer AUROCs. Cosine and JSD lenses agree on the U-shape but differ on which tap is sharpest: JSD concentrates discriminative mass at the canonical tap $\ell = 29$; cosine spreads it across many taps, producing a much larger vector-vs-best-single-layer gap.

Layer	Block type	ΔD_{jsd} AUROC	ΔD_{cos} AUROC
0	embed	0.605	0.519
7	`hcs` (shallow)	0.662	0.595
12	`hcm`	0.656	0.555
17	`attn`	0.723	0.646
24	`attn`	0.685	0.612
28	`hcs`	0.565	0.698
29 (canonical tap)	`hcm`	**0.794** (best JSD)	0.604
30 (post-norm−1)	—	0.512	**0.729** (best cos)
31 (idle, see App. A.1)	—	0.499 (degen.)	0.729 ($= h_{30}$)
32-d vector	all	0.823	0.844

E. Largest Q2 region detail and splice fine-profile

The largest single contiguous Q2 region on chr22 spans coordinates chr22:22,893,870–22,895,351 (1,481 bp). Mean settling depth within the region is $\bar{c} = 31.31$; mean PhyloP-100way is -0.63. The region overlaps an annotated intron of an inactivated gene model and is flanked by an ENCODE cCRE-ELS element 460 bp downstream. No coding sequence intersects the region. We list this single region as a representative case; 5,089 additional Q2 regions ≥ 100 bp are released as a BED file at `q2_regions.bed` in the supplementary materials.

On chr17 the donor profile reaches its mean settling-depth minimum at position $+20$ bp ($\bar{c} = 24.06$) and remains ≥ 1.5 layers below intronic baseline ($\bar{c} = 27.69$) out to ± 200 bp. On the same chromosome the acceptor profile reaches its minimum at $+50$ bp ($\bar{c} = 23.33$) and similarly remains depressed beyond ± 200 bp. On chr22 both donor and acceptor minima sit at coordinate 0 ($\bar{c} = 23.65$ and 23.64 respectively). The asymmetry — donor minimum exonic-side, acceptor minimum intronic-side — mirrors the asymmetric splice-grammar features (branch point, polypyrimidine tract) that lie predominantly on the intronic side of acceptors.

F. Detailed specifications of the four genomic foundation models

Table 7. Architectural specifications, tokenisation, per-window context, runtime, and per-model q_{70}-calibrated $\gamma_{\cos}$ for the four genomic foundation models compared in §3.3 and Appendix B.

Model	Architecture	Layers	Hidden	Tokens / 6 kb window	Wall-clock	$\gamma_{q_{70}}$
Evo 2 7B	Hybrid Transformer + StripedHyena 2	32	4096	6,000 (1 bp)	reused	0.396
HyenaDNA-large-1m	Pure Hyena	8	256	6,001 (1 bp + BOS)	~ 4 min	0.358
NT-v2 500M	Transformer MLM (k-mer)	29	1024	671 ($k = 6$, ~ 4 kb)	~ 7.5 min	0.533
DNABERT-2 117M	Transformer MLM (BPE)	12	768	~ 600 (BPE, ~ 3 kb)	~ 3 min	0.677

Table 8. The 32-d trajectory dominates any single-layer feature. The primary cosine lens reaches 0.844 with 32 layers and only 0.729 at its best single tap – a $+0.115$ gap. Bracketed numbers are $1{,}000$-bootstrap 95% confidence intervals.

Feature	dim.	Stratified 10-fold AUROC	LOGO-CV AUROC
Best single-layer $\Delta D_{\cos}$ ($\ell = 30$ tap)	1	0.729 $[0.717, 0.741]$	0.726 $[0.694, 0.758]$
Best single-layer ΔD_{jsd} ($\ell = 29$ canonical tap)	1	0.794 $[0.781, 0.806]$	0.787 $[0.752, 0.821]$
Evo 2 log-likelihood	1	0.751 $[0.738, 0.764]$	0.793 $[0.740, 0.846]$
32-d ΔD_{jsd} vector	32	0.823 $[0.813, 0.832]$	0.821 $[0.790, 0.853]$
32-d $\Delta D_{\cos}$ vector (primary)	32	$\mathbf{0.844}$ $[0.831, 0.857]$	$\mathbf{0.843}$ $[0.811, 0.876]$
Ensemble ($\Delta D +$ Evo 2 LL)	33	0.861 $[0.851, 0.871]$	0.866 $[0.832, 0.899]$

G. Variant pathogenicity AUROC – detailed panels

H. Canonical vs. non-canonical splice motif breakdown

A finer dissection of the splice axis by motif class – using the genomic dinucleotides immediately downstream of the donor and upstream of the acceptor – reveals an unexpected ordering: *non-canonical* splice donors (AT-AC and other rare motifs, pooled $\bar{c} = 25.13$, $n = 5{,}977$) converge *earlier* than the dominant canonical GT-AG donors ($\bar{c} = 25.79$, $n = 350{,}552$), while the minor canonical GC-AG class is the deepest of the three ($\bar{c} = 27.01$, $n = 5{,}165$). The universality vs. intron ($\bar{c} = 27.72$) holds for every splice class, but the within-class ordering is the opposite of a naive "stronger motif $\Rightarrow$ shallower recognition" prior. The Cohen's d between canonical GT-AG and non-canonical donors is small (≈ 0.05), consistent with the constrained branch-point and polypyrimidine context that flanks non-canonical introns (Wang & Burge, 2008; Burge et al., 1998); we report the magnitudes honestly rather than over-interpret. The finding replicates qualitatively on the 8-layer HyenaDNA-large at depth-normalised scale.

I. Cross-architecture two-tier structure – figure

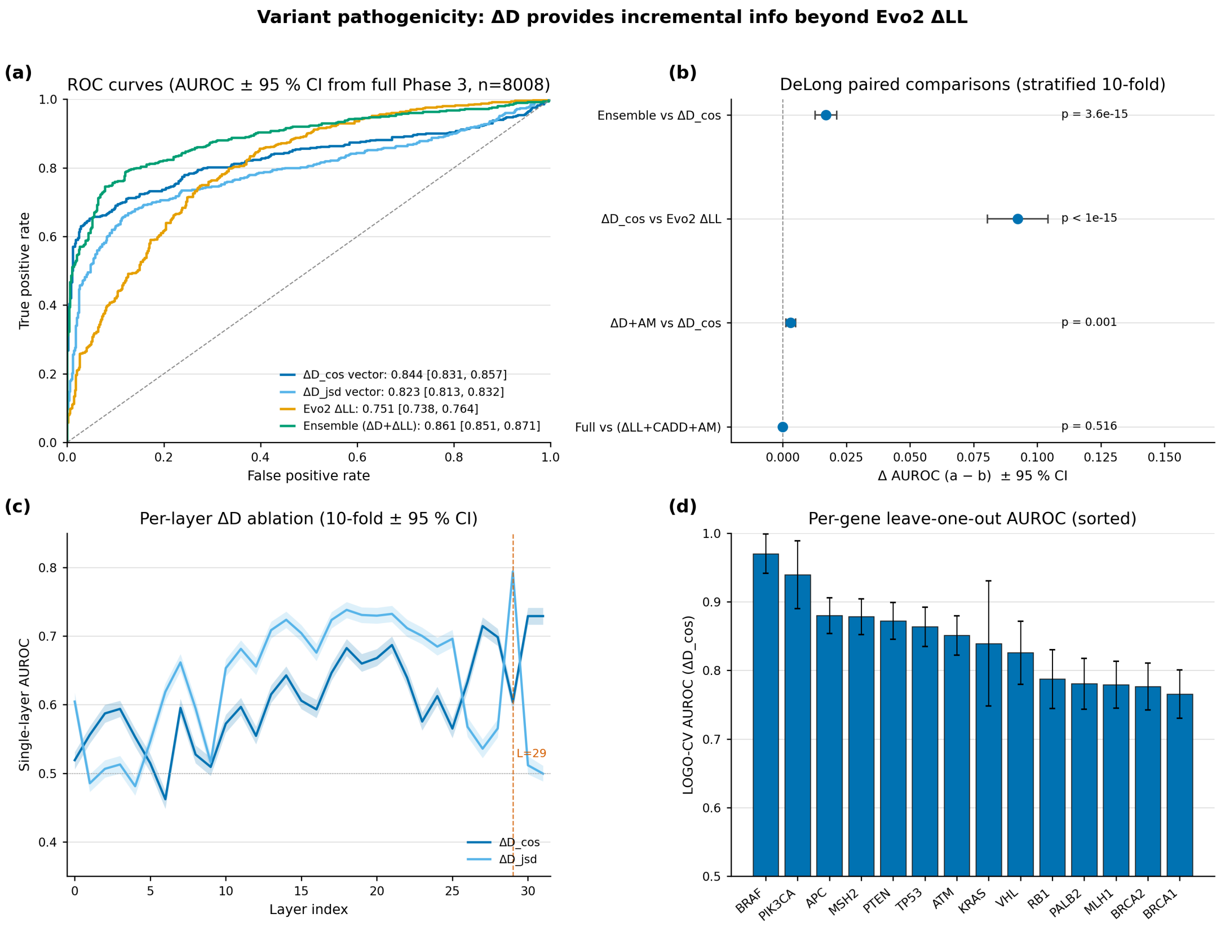

Figure 5. Variant pathogenicity discrimination on 8,008 ClinVar P/LP vs B/LB SNVs across 15 cancer-associated genes (the AUROC numerical summary is Table 8 above). **(a)** ROC curves for the four scoring features. **(b)** DeLong paired comparisons: ΔD adds significant information beyond Evo 2 ΔLL ($p = 3.6 \times 10^{-15}$). **(c)** Per-layer single-feature AUROC across all 32 layers – best single tap is $\ell = 29$ for JSD; the 32-d vector beats it by $+0.05$ to $+0.12$. **(d)** Leave-one-gene-out AUROC across 14 evaluable genes is uniformly high (≥ 0.77).

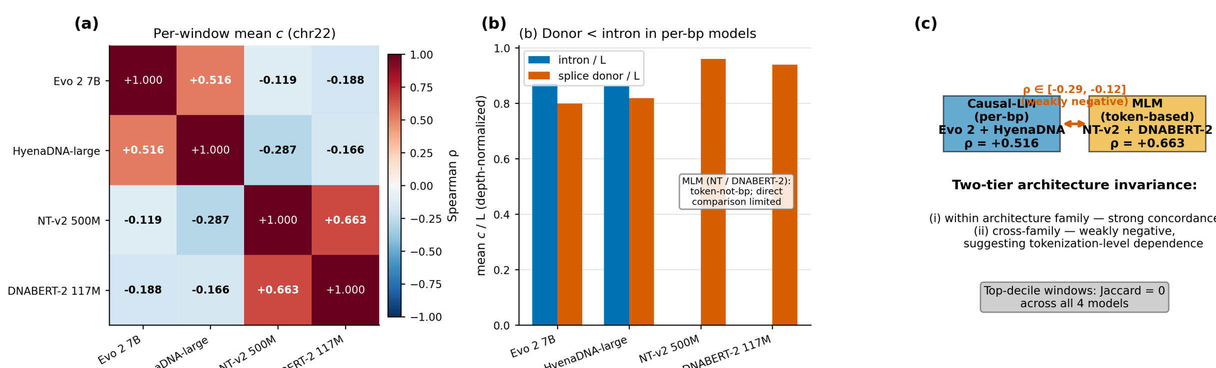

Figure 6. Cross-architecture two-tier structure (numerical detail in §3.3 body and Tables 4, 5). **(a)** Pairwise Spearman ρ of per-window mean settling depth across four genomic foundation models on chr22. **(b)** Depth-normalised donor vs intron mean c/L: per-bp causal-LMs (Evo 2, HyenaDNA) preserve donor $<$ intron; tokenised MLMs (NT-v2, DNABERT-2) tokenise at window level and therefore show no per-position separation. **(c)** Schematic of the two-tier structure with the within/cross-family ρ values.