



(A) Training loss



(B) Per-layer M2: baseline vs tuned



(C) SVD spectrum

